


HOG descriptor




HOG

Dalal, Triggs. **Histograms of Oriented Gradients** for Human Detection. CVPR, 2005




HOG

Dalal, Triggs. **Histograms of Oriented Gradients** for Human Detection. CVPR, 2005



histogram of
'unsigned' gradients




soft binning



gradient magnitude
(one for each cell)


HOG

Dalal, Triggs. **Histograms of Oriented Gradients** for Human Detection. CVPR, 2005



Block
(2x2 cells)

histogram of
'unsigned' gradients



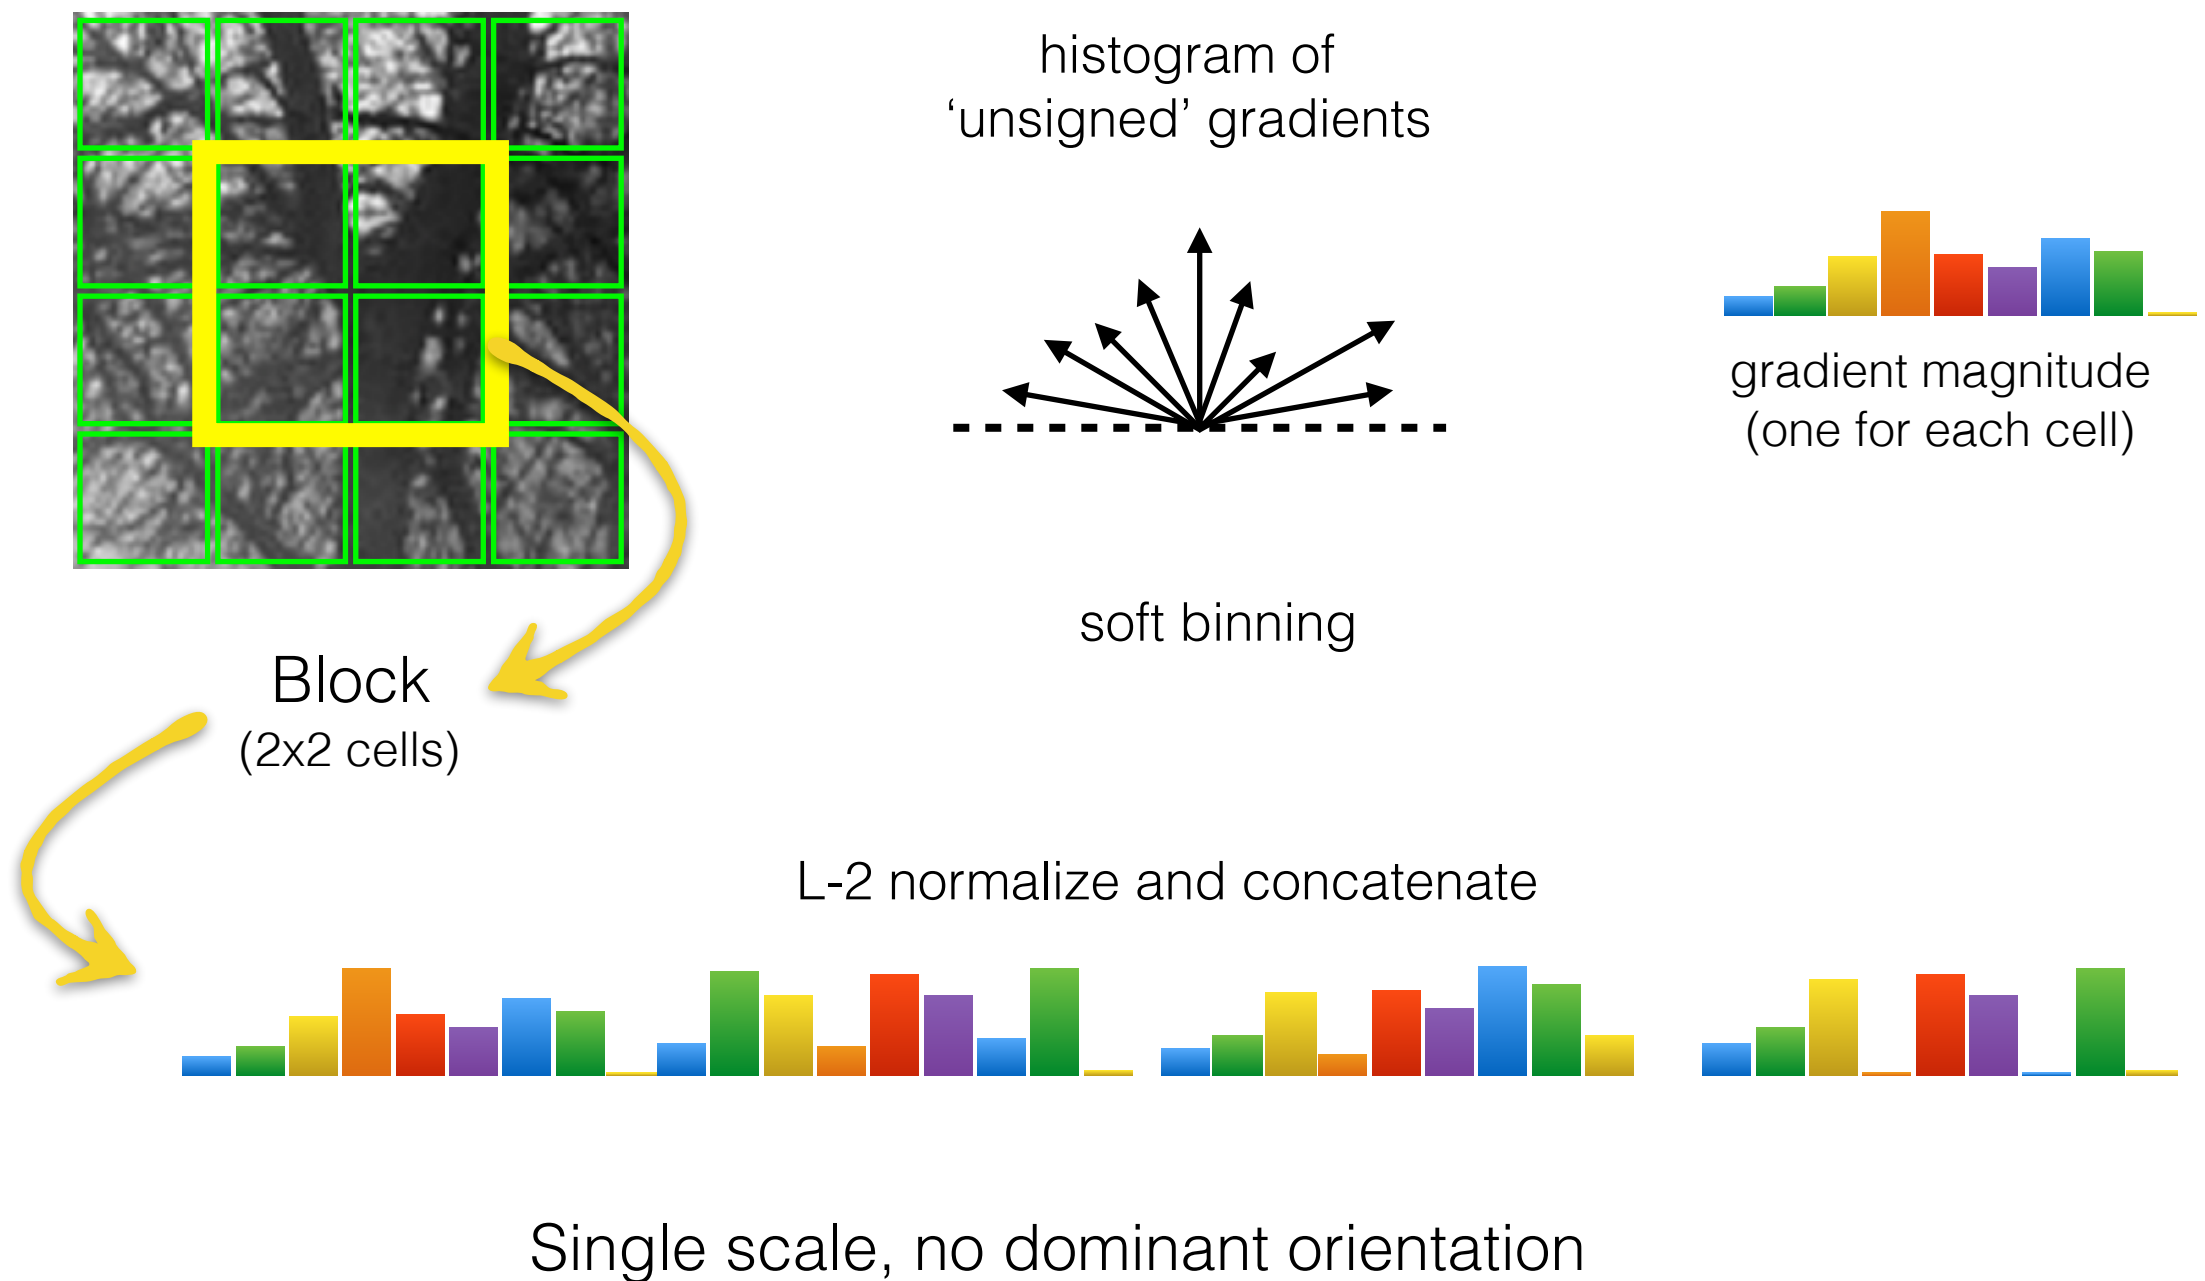
soft binning



gradient magnitude
(one for each cell)

HOG

Dalal, Triggs. **Histograms of Oriented Gradients** for Human Detection. CVPR, 2005



Pedestrian detection


1 cell step size



128 pixels
16 cells
15 blocks

$$15 \times 7 \times 4 \times 36 = 3780$$

visualization



64 pixels
8 cells
7 blocks

Redundant representation due to overlapping blocks

How many times is each inner cell encoded?

