

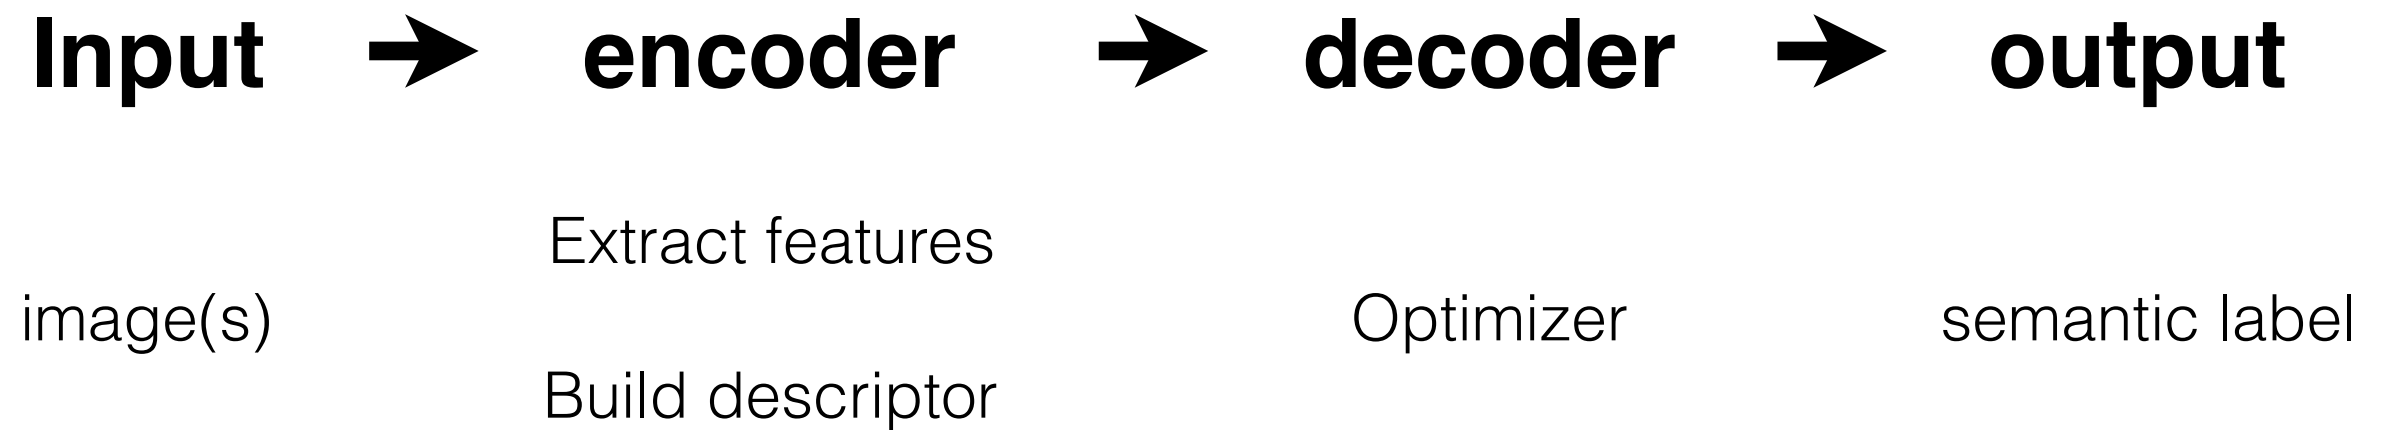
BoW Classification

Computer Vision

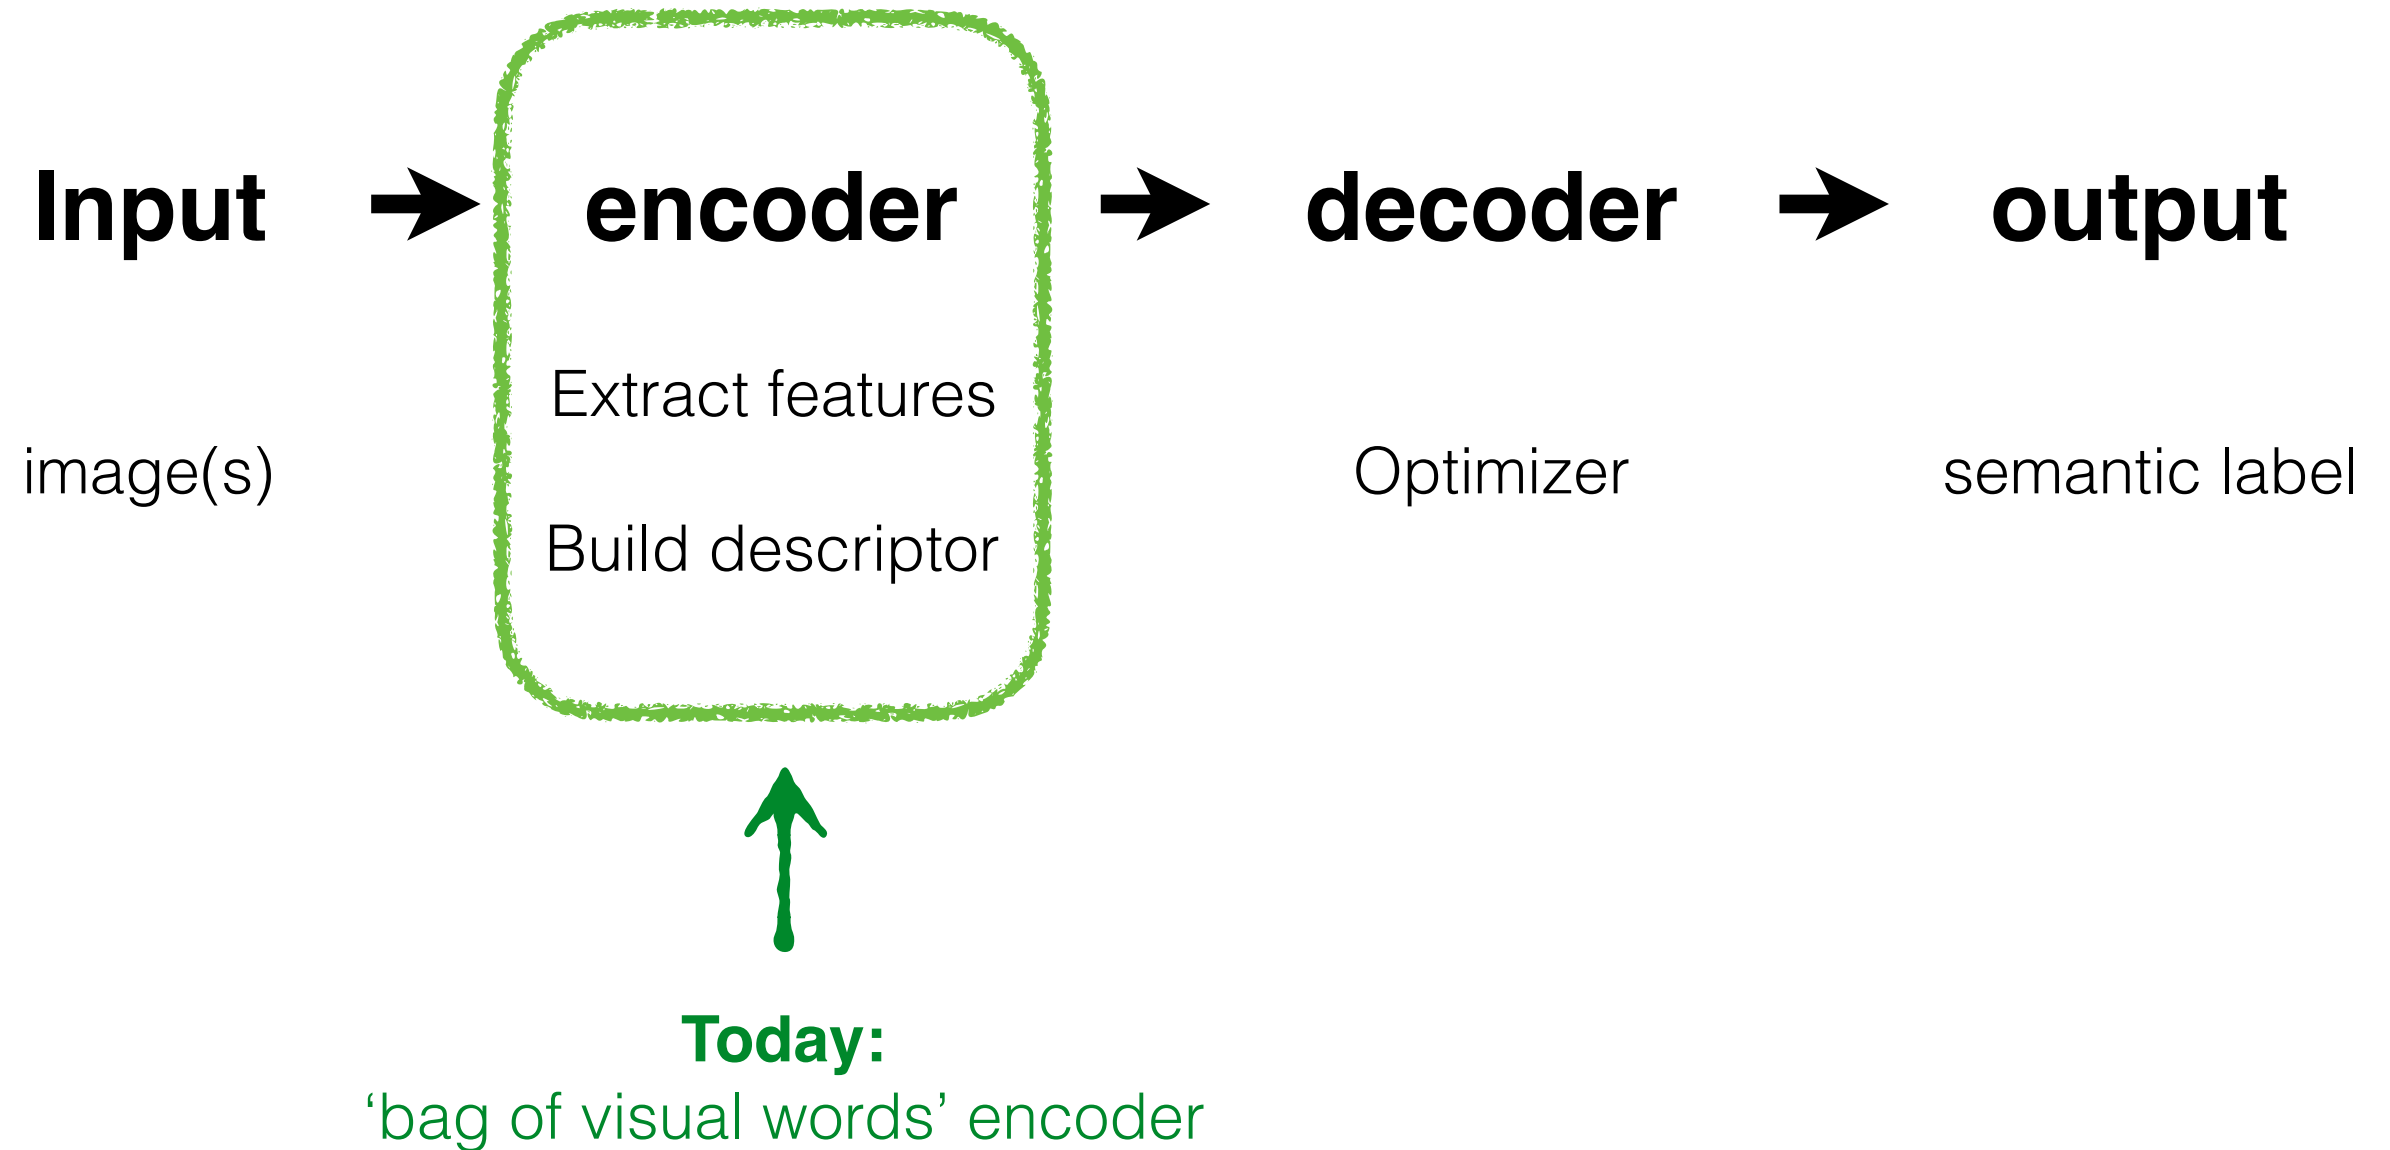
Carnegie Mellon University (Kris Kitani)

‘Classical’

Image Classification Pipeline



^{'Classical'} Image Classification Pipeline



BoW encoder

Encode:


1. Dictionary Learning
2. Build Bags-of-Words (BOW) descriptor

Decode:

1. Train classifier
2. Classify using BOW descriptor

BoW encoder

a. Extract features from image



Encode:

1. Dictionary Learning

2. Build Bags-of-Words
(BOW) descriptor


Decode:

1. Train classifier

2. Classify using BOW
descriptor

BoW encoder

a. Extract features from image



Encode:

1. Dictionary Learning

2. Build Bags-of-Words
(BOW) descriptor

What kinds of features can we extract?



Decode:

1. Train classifier

2. Classify using BOW
descriptor

BoW encoder

a. Extract features from image



What kinds of features can we extract?

Encode:

1. Dictionary Learning

2. Build Bags-of-Words
(BOW) descriptor

Decode:

1. Train classifier

2. Classify using BOW
descriptor

Regular grid

- Vogel & Schiele, 2003
- Fei-Fei & Perona, 2005

Interest point detector

- Csurka et al. 2004
- Fei-Fei & Perona, 2005
- Sivic et al. 2005

Other methods


- Random sampling (Vidal-Naquet & Ullman, 2002)
- Segmentation-based patches (Barnard et al. 2003)

BoW encoder


a. Extract features from image



b. Unsupervised Clustering



K-means clustering



Encode:

1. Dictionary Learning


2. Build Bags-of-Words
(BOW) descriptor

Decode:



1. Train classifier

2. Classify using BOW
descriptor

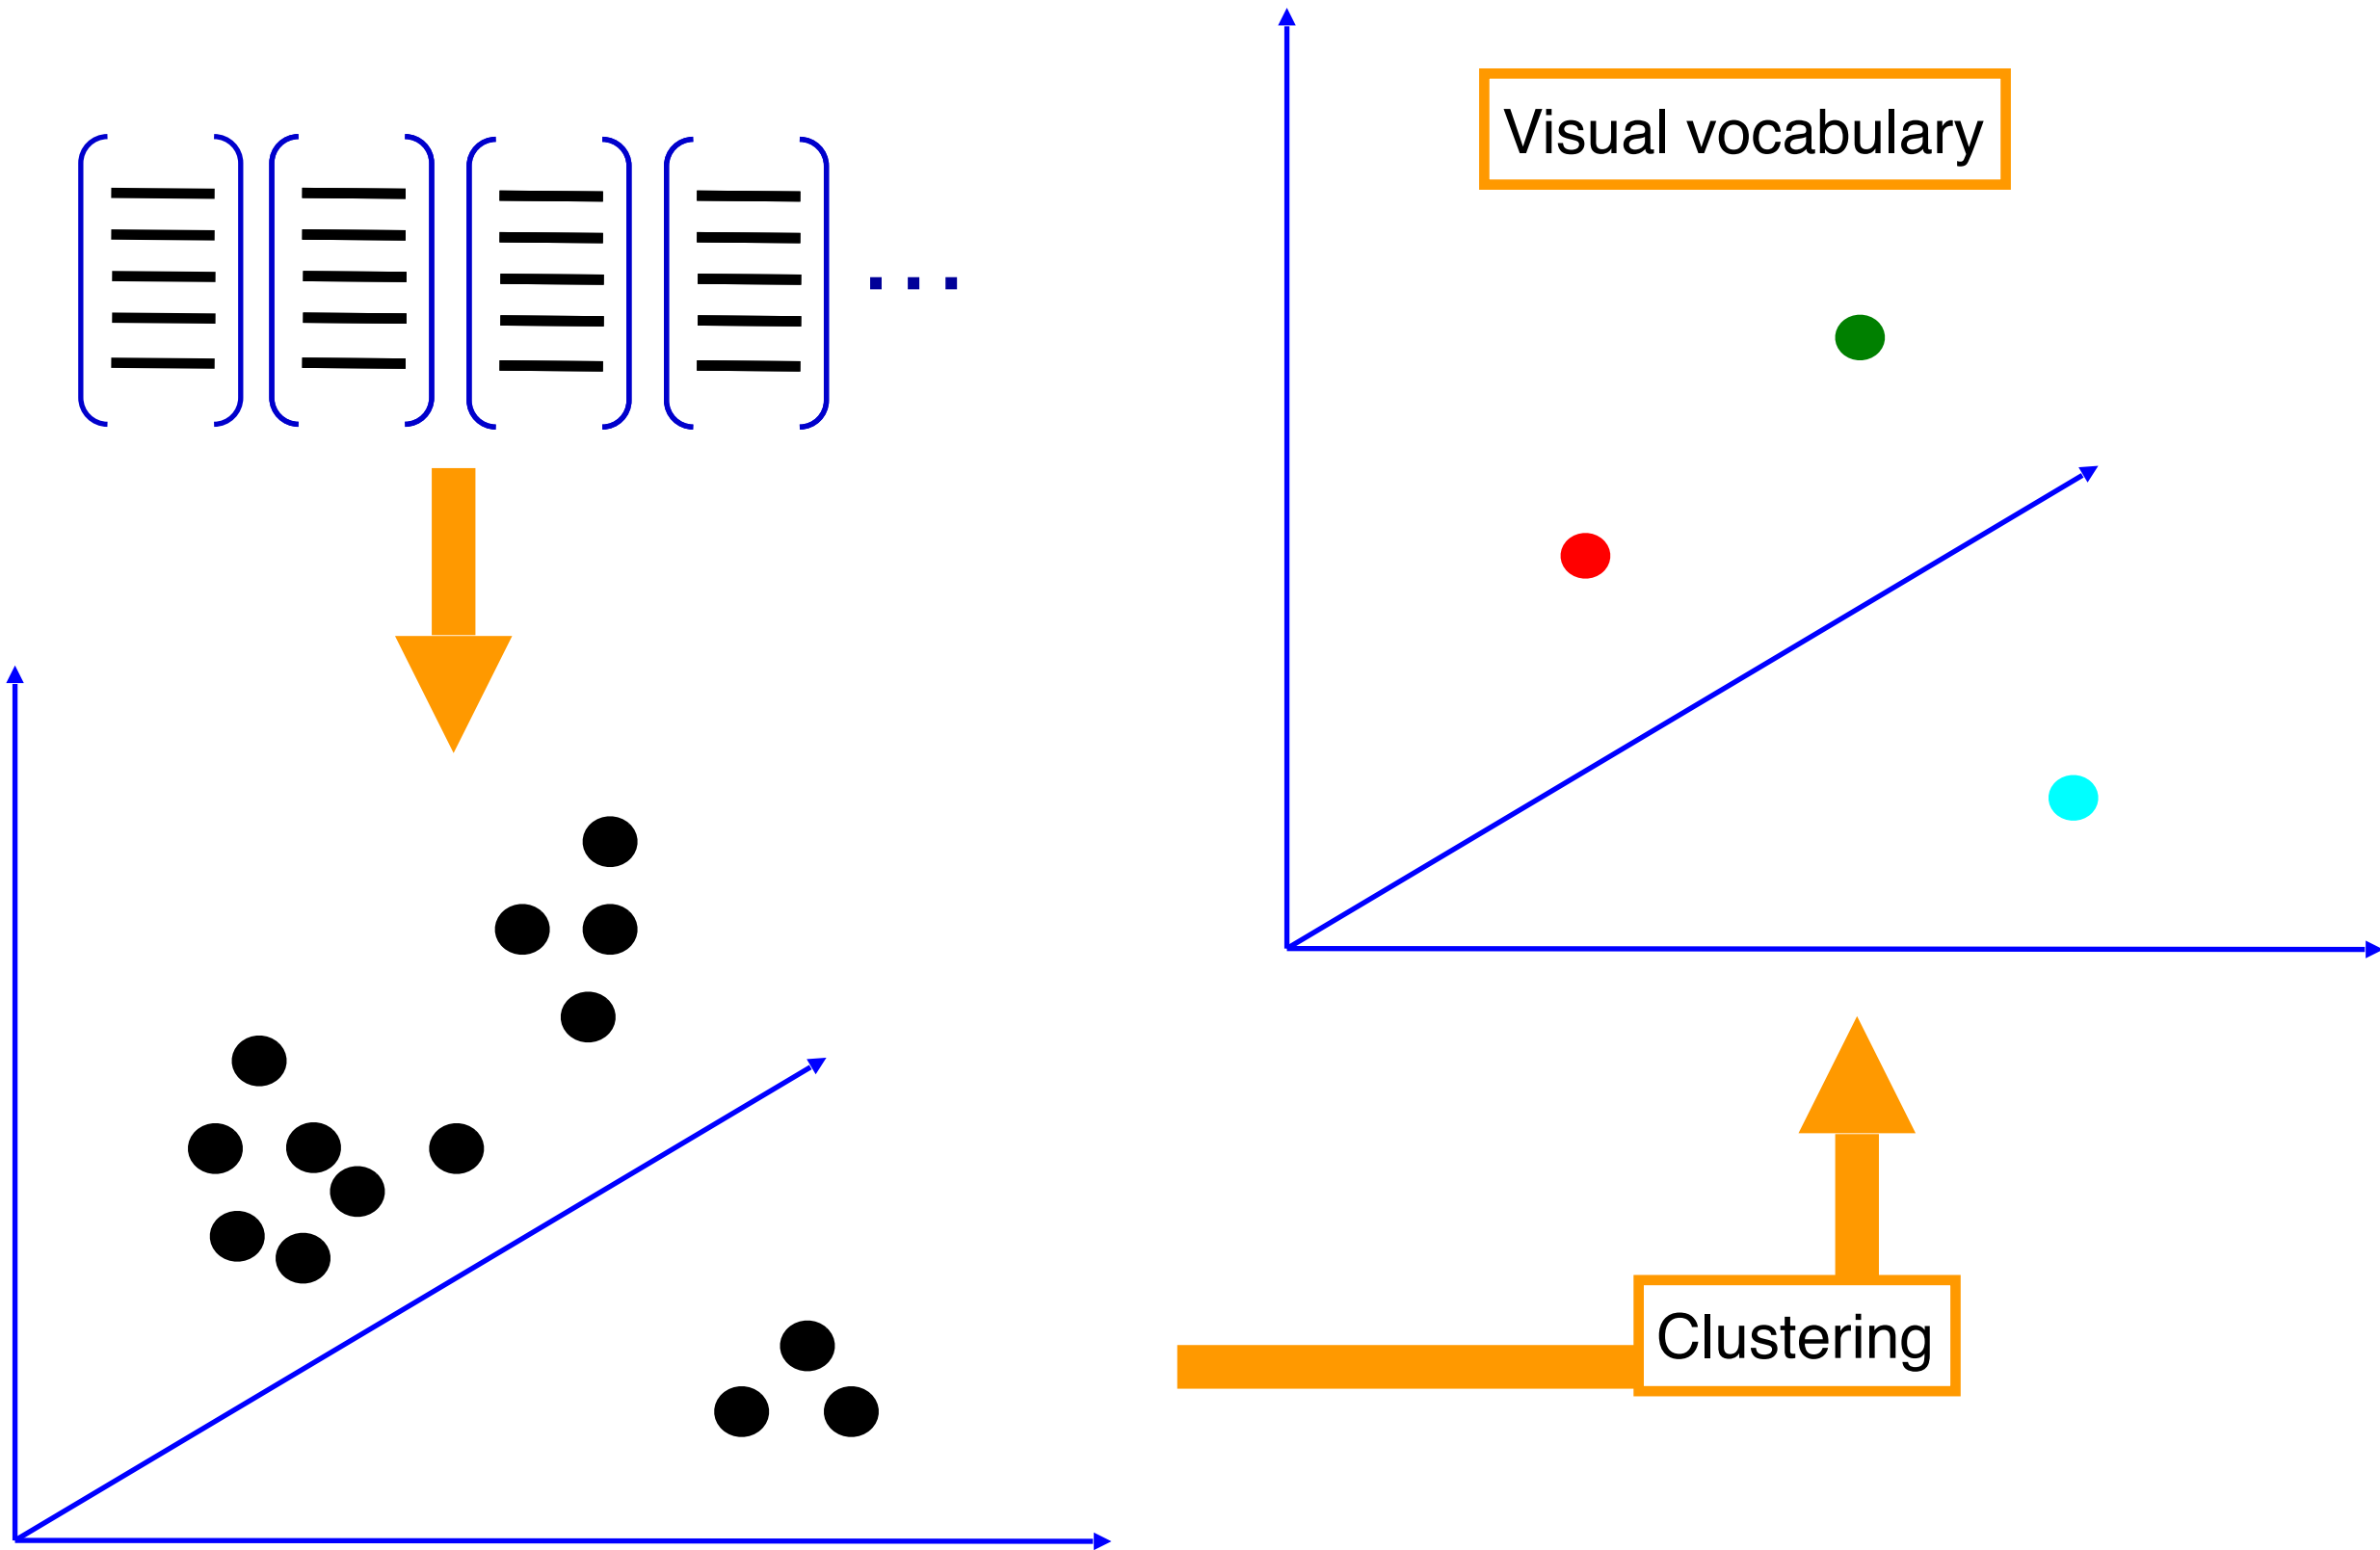
b. Unsupervised Clustering




b. Unsupervised Clustering



b. Unsupervised Clustering




b. Unsupervised Clustering




1. Select initial
centroids at random

b. Unsupervised Clustering




1. Select initial centroids at random




2. Assign each object to the cluster with the nearest centroid.



b. Unsupervised Clustering



1. Select initial centroids at random




2. Assign each object to the cluster with the nearest centroid.




3. Compute each centroid as the mean of the objects assigned to it (go to 2)


b. Unsupervised Clustering




1. Select initial centroids at random



2. Assign each object to the cluster with the nearest centroid.




3. Compute each centroid as the mean of the objects assigned to it (go to 2)




2. Assign each object to the cluster with the nearest centroid.


b. Unsupervised Clustering





1. Select initial centroids at random



2. Assign each object to the cluster with the nearest centroid.



3. Compute each centroid as the mean of the objects assigned to it (go to 2)



2. Assign each object to the cluster with the nearest centroid.

Repeat previous 2 steps until no change

b. Unsupervised Clustering

K-means Clustering

Given k :

1. Select initial centroids at random.
2. Assign each object to the cluster with the nearest centroid.
3. Compute each centroid as the mean of the objects assigned to it.
4. Repeat previous 2 steps until no change.

BoW encoder

Encode:

1. Dictionary Learning

2. Build Bags-of-Words
(BOW) descriptor



After learning a visual dictionary,
we need an image representation.

Let's use a bag-of-words descriptor!

Decode:

1. Train classifier

2. Classify using BOW
descriptor

BoW encoder

a. Quantize feature



Encode:

1. Dictionary Learning
2. Build Bags-of-Words (BOW) descriptor

Decode:

1. Train classifier
2. Classify using BOW descriptor

BoW encoder

Encode:

1. Dictionary Learning
2. Build Bags-of-Words (BOW) descriptor


Decode:

1. Train classifier
2. Classify using BOW descriptor

a. Quantize feature



Associate each feature to nearest 'word' (cluster)



BoW encoder

Encode:

1. Dictionary Learning
2. Build Bags-of-Words (BOW) descriptor


Decode:

1. Train classifier
2. Classify using BOW descriptor


a. Quantize feature



Associate each feature to nearest 'word' (cluster)



b. Build histogram (for each image)



BoW encoder

Encode:

1. Dictionary Learning
2. Build Bags-of-Words (BOW) descriptor


Decode:

1. Train classifier
2. Classify using BOW descriptor

BoW encoder

Encode:


1. Dictionary Learning
2. Build Bags-of-Words (BOW) descriptor



Given the bag-of-features representations of images from different classes...

Decode:

1. Train classifier
2. Classify using BOW descriptor



Train some fancy machine learning model


BoW encoder

Encode:

1. Dictionary Learning
2. Build Bags-of-Words (BOW) descriptor

Decode:

1. Train classifier
2. Classify BOW descriptor



Get output from
machine learning
model

^{'Classical'} Image Classification Pipeline

