

**Generative Adversarial Networks**  
Computer Vision  
**Carnegie Mellon University (Kris Kitani)**

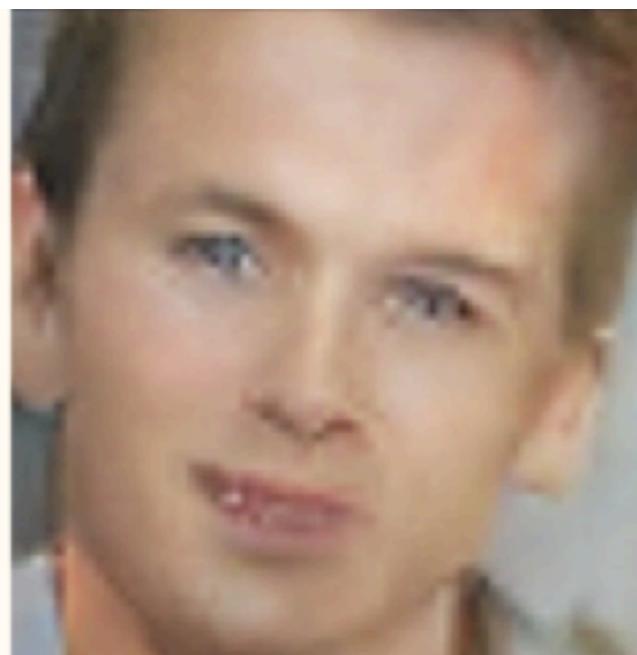
# Progress of Face Image Generation

[Brundage et al 2018]



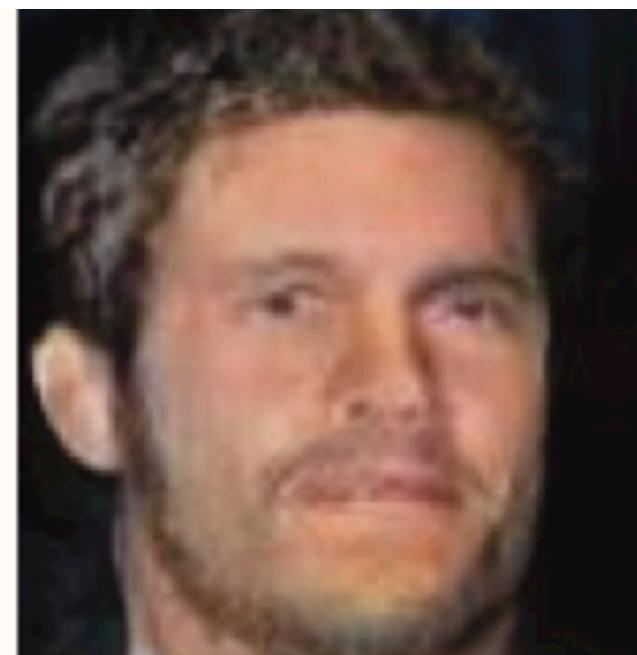
2014

[Goodfellow et al.]



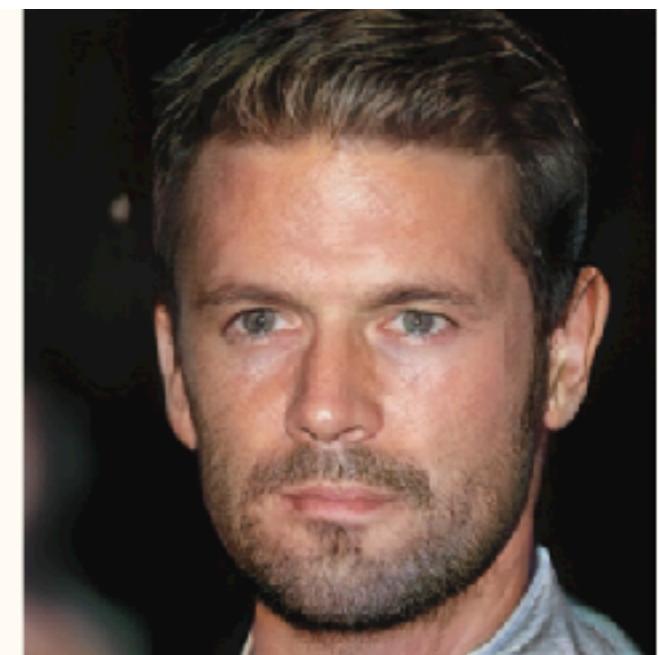
2015

[Radford et al.]



2016

[Liu & Tunzel]



2017

[Karras et al.]



**BigGAN [Brock et al. 18]**



**[Jin et al. 17]**



**[Wang et al. 17]**

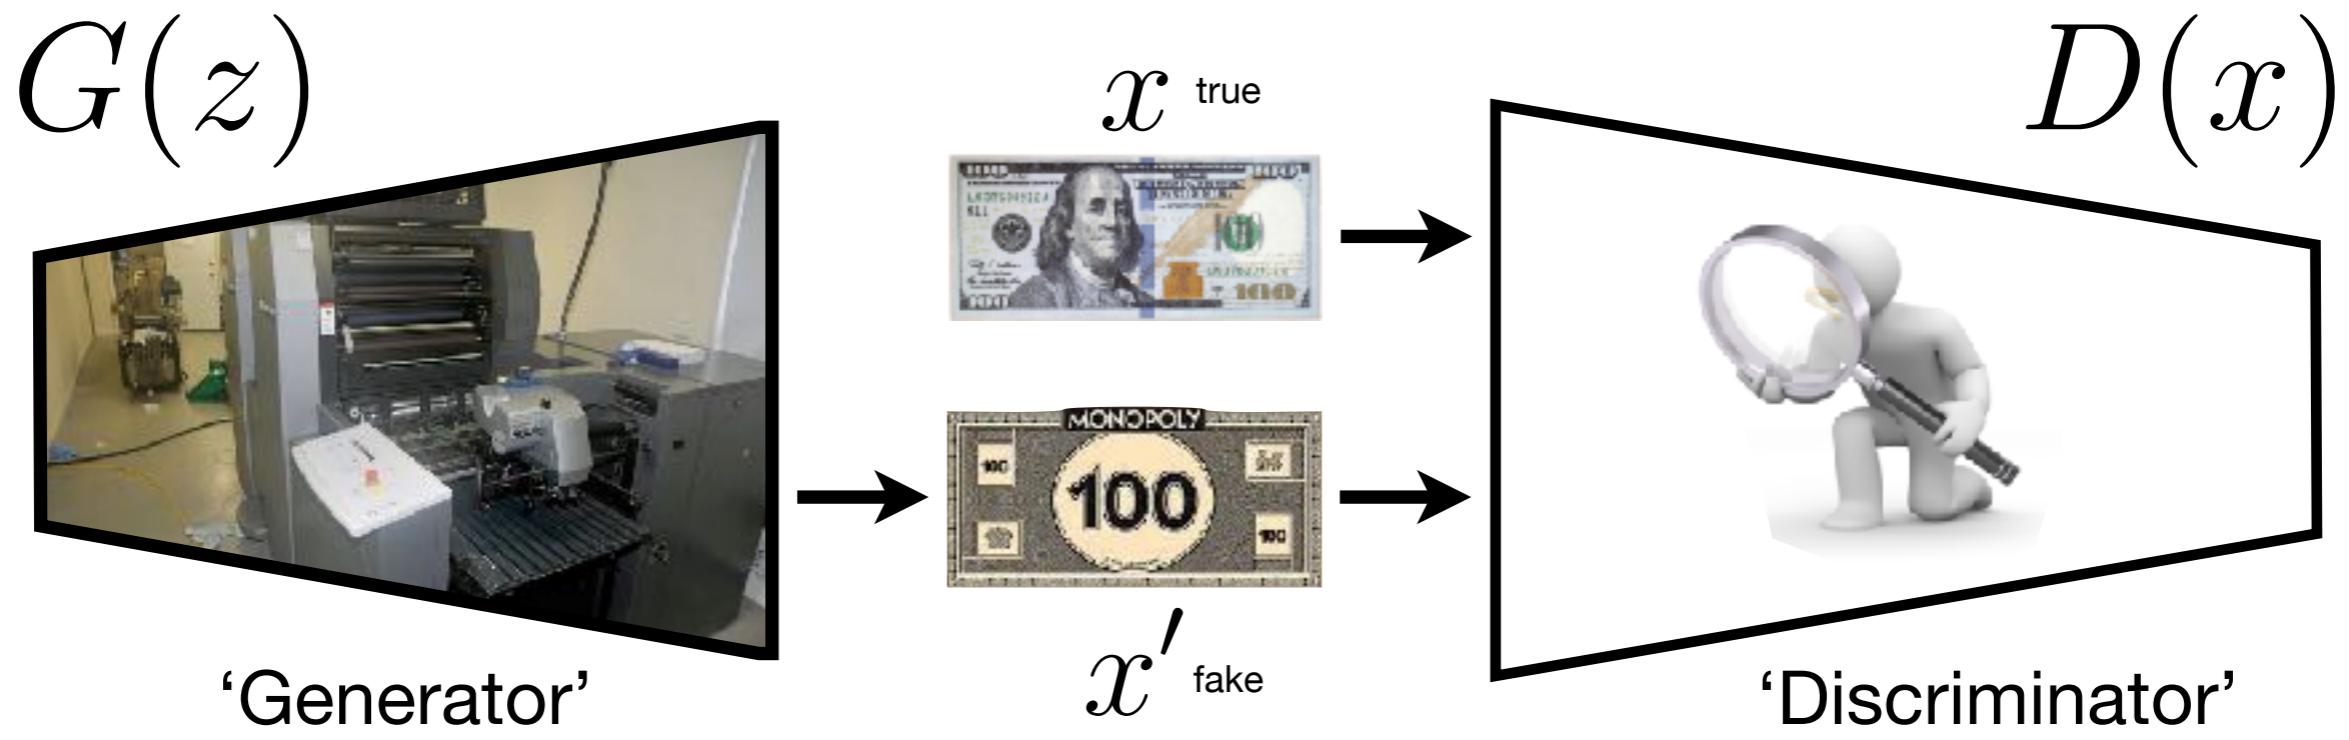


**[Taigaman et al. 16]**

**CycleGAN [Zhu et al. 17]**

Counterfeiters trying to generate fake money

Police trying to detect fake currency



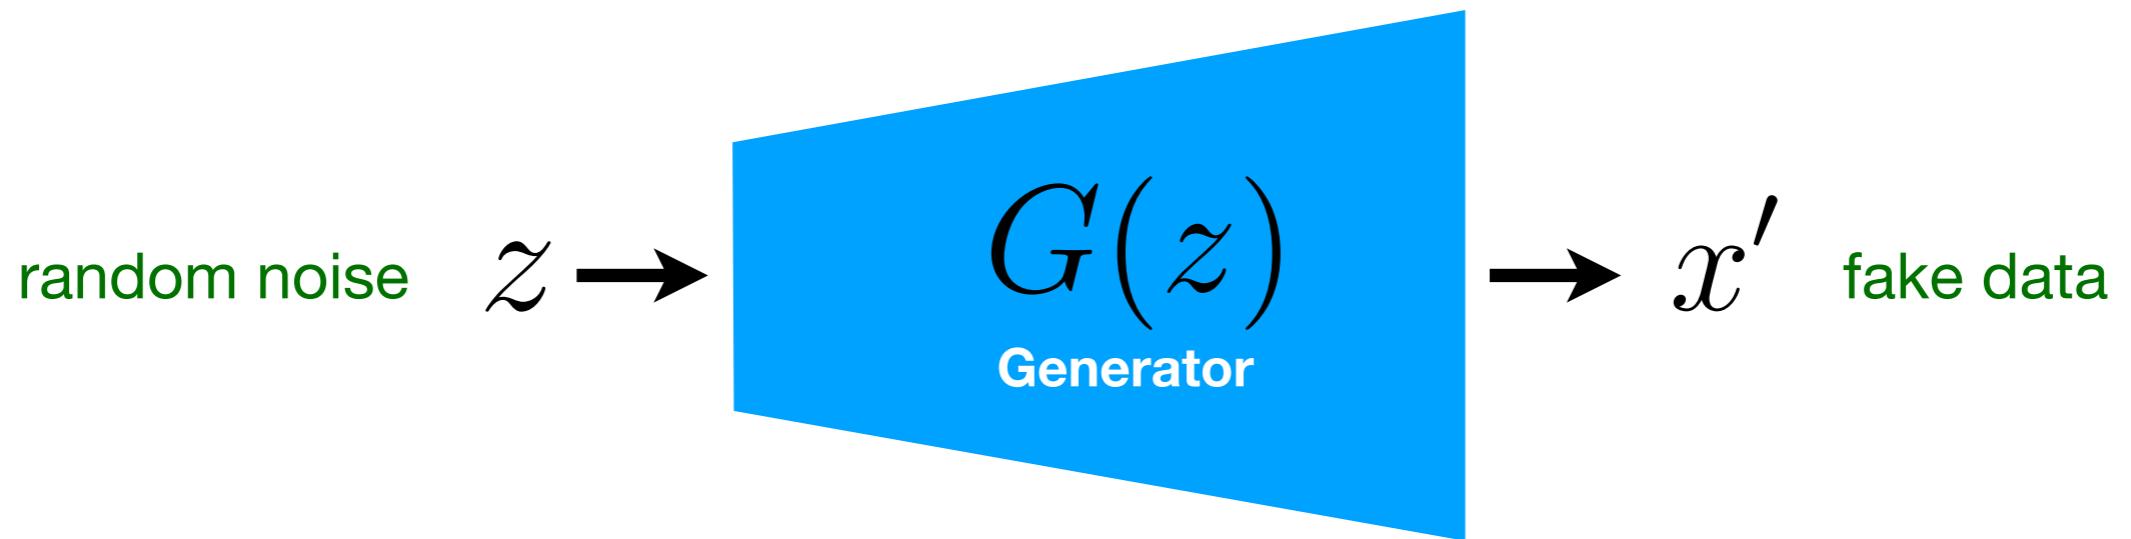
**Counterfeiters get better**  
to avoid getting caught

**Police get better**  
at identifying fakes

# Motivating the Discriminator

# How would you score a Generator?

Assume someone gives you a generator:



Generator takes noise and transforms it to a data point (e.g., image, face, trajectory)

$$z \sim q(z)$$

$$x' \sim G(z)$$

If the generator is very good, the fake data should look really 'similar' to the true data

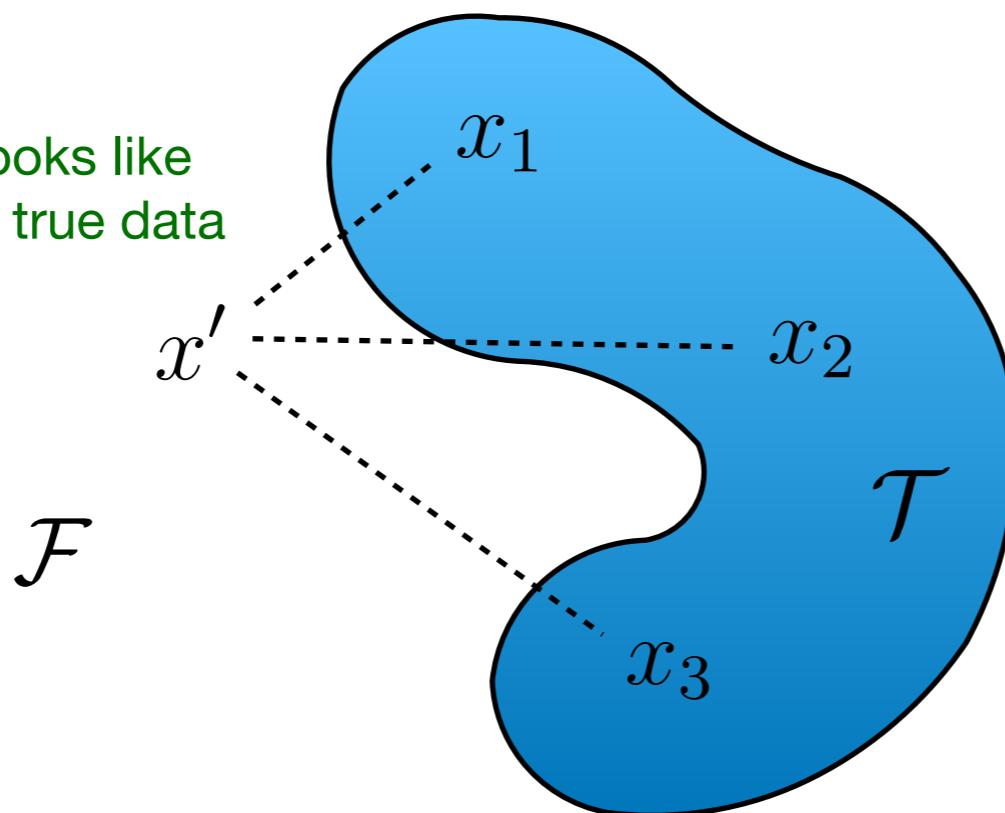
# How would you score a Generator?

If the generator is very good, the fake data should look really ‘similar’ to the true data

*How should we score a generator?*

We could **define** some **distance function**...

Heuristic: Fake data looks like true data if it is close to true data



But what does it mean for a data point to look similar to true data?

*But we could also learn it ...*

# Learning a Distance Function

Given:

$$\mathcal{T} = \{x_1, \dots, x_N\} \quad \mathcal{F} = \{x'_1, \dots, x'_N\}$$

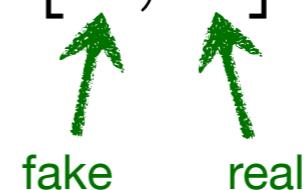
True data

Fake data

We want to learn:

$$D(x; \theta_D) \in [0, 1]$$

(inverse) distance function



***How would you learn such a distance function?***

(supervised learning)

# Learning a Distance Function

Given:

$$\mathcal{T} = \{x_1, \dots, x_N\} \quad \mathcal{F} = \{x'_1, \dots, x'_N\}$$

True data

Fake data

We want to learn:

$$D_\theta(x) = \begin{cases} 1 & x \text{ is true} \\ 0 & \text{otherwise} \end{cases}$$

Solve this:

$$\max_{\theta} \left\{ \sum_{x \in \mathcal{T}} [D_\theta(x)] + \sum_{x' \in \mathcal{F}} [1 - D_\theta(x')] \right\}$$

For all true data                                  For all fake data

Should be 1 for true data

Should be 1 for fake data...

since this term  
should be zero

If  $D$  is differentiable, it can be learned with gradient descent.

# Learning a Discriminator Function

Given:

$$\mathcal{T} = \{x_1, \dots, x_N\} \quad \mathcal{F} = \{x'_1, \dots, x'_N\}$$

True data

Fake data

We want to learn:

$$D_\theta(x) = \begin{cases} 1 & x \text{ is true} \\ 0 & \text{otherwise} \end{cases}$$

Solve this:

$$\max_{\theta} \left\{ \sum_{x \in \mathcal{T}} [D_\theta(x)] + \sum_{x' \in \mathcal{F}} [1 - D_\theta(x')] \right\}$$

For all true data                                  For all fake data

Should be 1 for true data

Should be 1 for fake data...

since this term  
should be zero

Let's call this the '**Discriminator**'

Because it differentiates between true and false data

If we had a generator and some true data,  
now we know how to train a discriminator

$$\max_{\theta} \left\{ \sum_{x \in \mathcal{T}} [D_{\theta}(x)] + \sum_{x' \in \mathcal{F}} [1 - D_{\theta}(x')] \right\}$$

$$\max_{\theta} \left\{ \sum_{x \sim p(x)} [D_{\theta}(x)] + \sum_{z \sim q(z)} [1 - D_{\theta}(G(z))] \right\}$$

alternative notation (explicit generator)

# Motivating the Generator

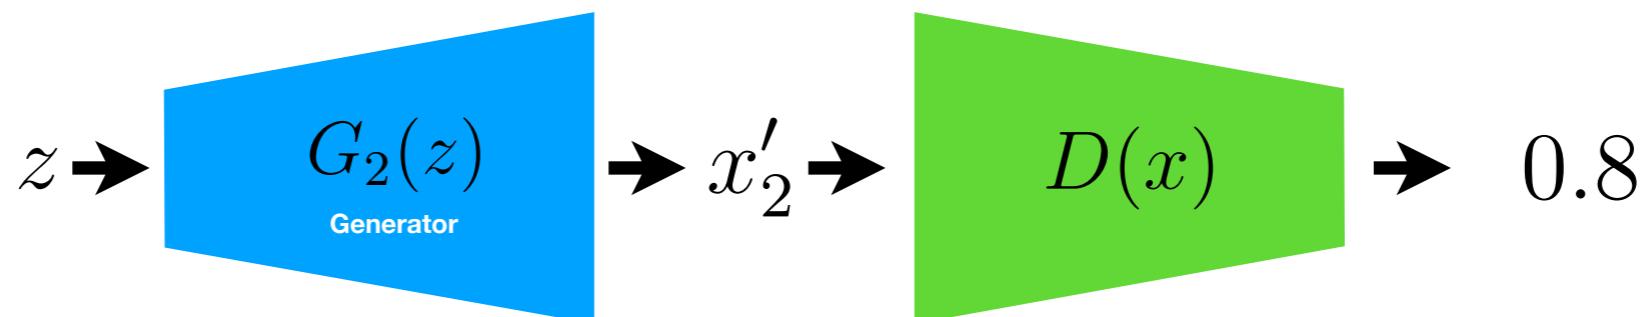
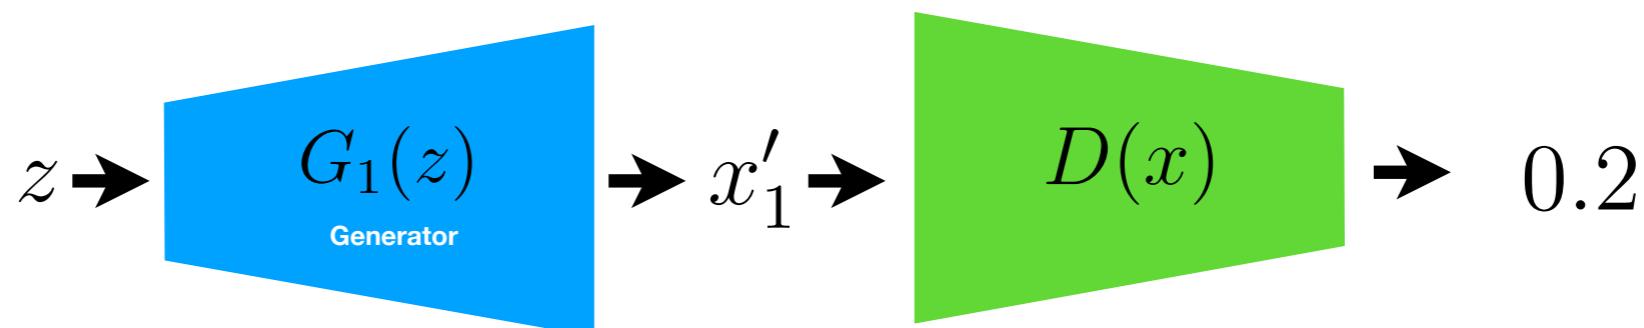
# How would you score a Generator?

Assume someone gives you a **discriminator**:

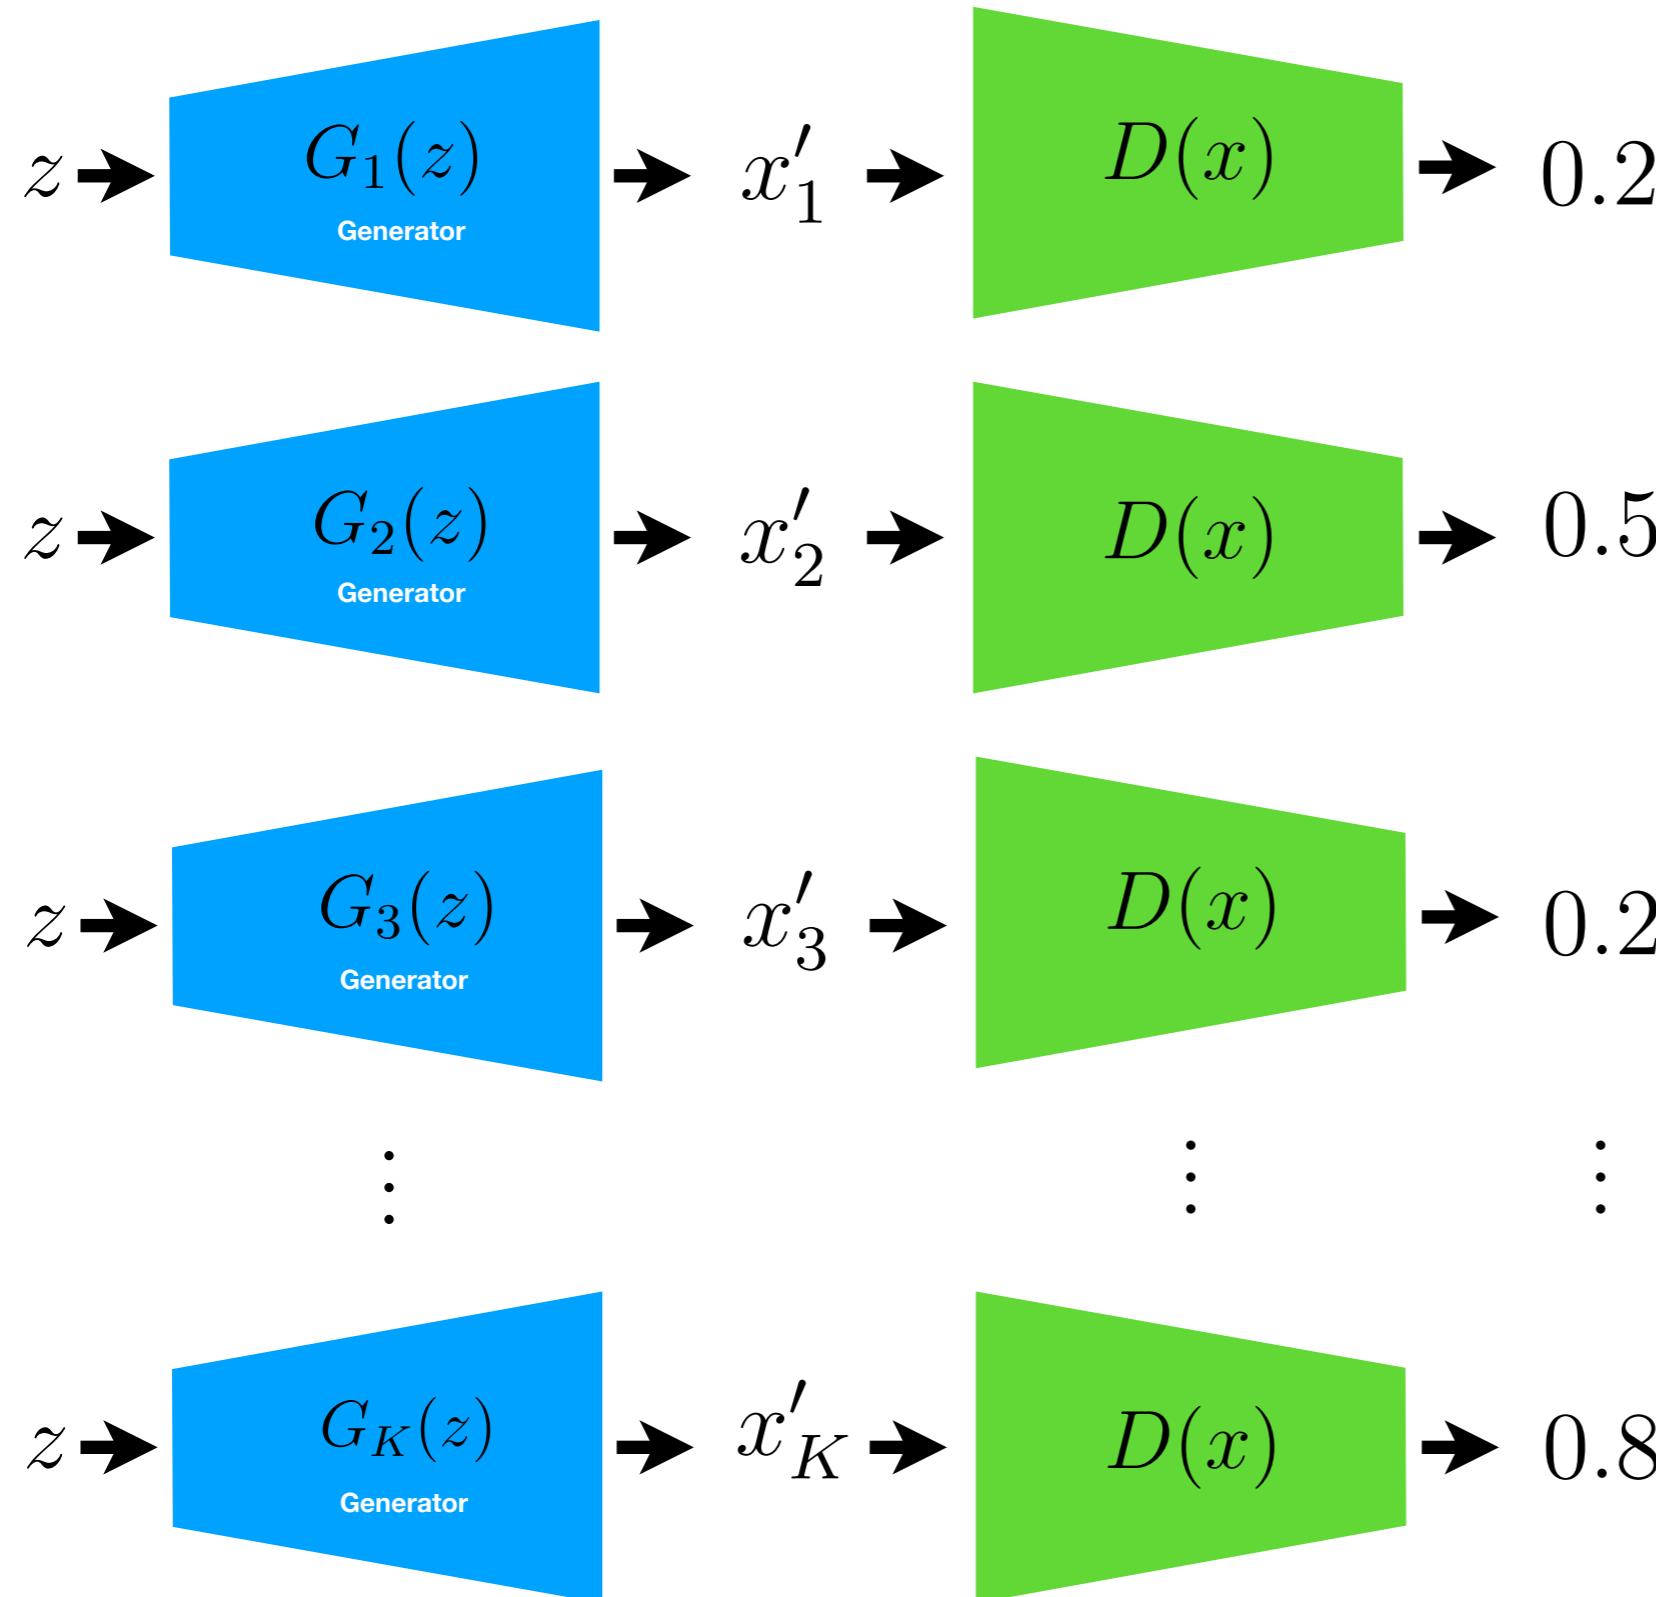
$$D_\theta(x) = \begin{cases} 1 & x \text{ is true} \\ 0 & \text{otherwise} \end{cases}$$

If you had a really good generator,  
it should be able to trick the discriminator

*Which generator is better? Why?*



# How would you find the best generator?



**Brute Force?**

You could get a whole bunch  
of random generators and  
keep the best one

But that could take a really  
long time!

We could also learn it ...

# Learning a Generator

Given:  $D_\theta(x) = \begin{cases} 1 & x \text{ is true} \\ 0 & \text{otherwise} \end{cases}$

***Could we learn a generator if we had access to only the discriminator function?***

(yes, supervised learning)

# Learning a Generator

Given:  $D_\theta(x) = \begin{cases} 1 & x \text{ is true} \\ 0 & \text{otherwise} \end{cases}$

***Could we learn a generator if we had access to only the discriminator function?***

$$\min_{\phi} \left\{ \sum_{z \sim p(z)} [1 - D(G_{\phi}(z))] \right\}$$

Should be close to 0 for fake data...  
since this term should be  
close to one

If  $G$  is differentiable, it can be learned with gradient descent.

If we had a discriminator,  
now we know how to train a generator

$$\min_{\phi} \left\{ \sum_{z \sim p(z)} [1 - D(G_{\phi}(z))] \right\}$$

Generator and true data → Discriminator

$$\max_{\theta} \left\{ \sum_{x \sim p(x)} [D_{\theta}(x)] + \sum_{z \sim q(z)} [1 - D_{\theta}(G(z))] \right\}$$

Discriminator → Generator

$$\min_{\phi} \left\{ \sum_{z \sim p(z)} [1 - D(G_{\phi}(z))] \right\}$$

*How do we put these two ideas together?*

# **Formulation**

*How do we put these two ideas together?*

$$\max_{\theta} \left\{ \sum_{x \sim p(x)} [D_{\theta}(x)] + \sum_{z \sim q(z)} [1 - D_{\theta}(G(z))] \right\}$$

Learn discriminator, given generator

$$\min_{\phi} \left\{ \sum_{z \sim p(z)} [1 - D(G_{\phi}(z))] \right\}$$

Learn generator, given discriminator

We learn both the generator and the discriminator at the same time!

$$\min_{\phi} \max_{\theta} \left\{ \sum_{\substack{x \sim p(x) \\ \text{true data}}} [D_{\theta}(x)] + \sum_{\substack{z \sim q(z) \\ \text{random noise}}} [1 - D_{\theta}(G_{\phi}(z))] \right\}$$

Should be 1  
Fake data  
Should be 0 (discriminator)  
Should be 1 (generator)

**Formulation in original paper...**

$$\min_{\phi} \max_{\theta} \sum_{x \sim p(x)} [\log D_{\theta}(x)] + \sum_{z \sim q(z)} [\log\{1 - D_{\theta}(G_{\phi}(z))\}]$$

```
1: function GAN( $\mathcal{T}$ )
2:   for  $t = 1, \dots, T$  do
3:     Noise  $\{z_m\}_{m=1}^M$  sampled from  $q(z)$ 
4:     True data  $\{x_m\}_{m=1}^M$  sampled from  $\mathcal{T}$ 
5:      $\phi = \phi + \nabla_\phi \frac{1}{M} \sum_m [\log D_\theta(x_m) + \log(1 - D_\theta(G_\phi(z_m)))]$ 
6:     Noise  $\{z_n\}_{n=1}^N$  sampled from  $q(z)$ 
7:      $\theta = \theta - \nabla_\theta \frac{1}{N} \sum_n \log(1 - D_\theta(G_\phi(z_n)))$ 
```

**Only 6 lines of code!**

(if you ignore the millions of gradient updates in line 5 and 7)